

# affirmoTM tutorial 3: settings

Each setting in settings.yaml, and how to change them

2026-10-04

## Table of contents

|          |                                    |          |
|----------|------------------------------------|----------|
| <b>1</b> | <b>The settings file</b>           | <b>1</b> |
| <b>2</b> | <b>Changing settings</b>           | <b>3</b> |
| 2.1      | Discount rate . . . . .            | 3        |
| 2.2      | Time horizon . . . . .             | 3        |
| 2.3      | Side by side . . . . .             | 3        |
| 2.4      | Willingness-to-pay lines . . . . . | 4        |
| 2.5      | Seed and number of runs . . . . .  | 4        |

This tutorial starts from a fresh analysis folder with the example input files (see tutorial 1):

```
library(affirmoTM)
```

```
tm_example_inputs()
```

## 1 The settings file

inputs/settings.yaml holds the run settings, one per line, with comments:

```
cat(readLines("inputs/settings.yaml"), sep = "\n")
```

```
# affirmoTM run settings.
```

```
# Stratifier columns in the transition matrices (and the cohort, entry states, and optionally costs and  
# utilities). Every column of a matrix that is not `from` and not listed here is read as a state.
```

```
# The age-group column must be called `agegrp`: it is the only stratifier that changes as patients age.
```

```
# Any other name (for example `age_band`) is read as a fixed stratifier, so patients would keep their
```

```
# entry age group for the whole horizon. Use `strata: []` for matrices without stratifiers.
```

```
strata: [agegrp, sex, mm_cluster]
```

```
# Seed for the random draws (entry states and yearly transitions), so a run can be repeated exactly.
```

```
seed: 2026
```

```
# Time horizon in cycles (years), counted from cycle 0 (2013). Everyone stops at the end of cycle
```

```
# horizon - 1, so patients who enter later are followed for fewer years.
```

```
horizon: 10
```

```
# A patient who would pass the oldest age group (for example reach 110 when the oldest group is
```

```
# 105-109) dies in that year, in this death state.
```

```

age_limit_death_state: CurrentDeath

# Annual discount rate for costs and QALYs (0.03 = 3%). Discounting runs from cycle 0 (2013) for
# everyone: a year in cycle t counts 1 / (1 + rate)^t.
discount_rate: 0.03

# Number of times the whole simulation is run. Run r uses seed + r - 1, so every run differs and the
# whole set can be repeated exactly. Each run gives one point on the cost-effectiveness plane.
n_runs: 100

# Events to count in the results. Each year a patient spends in one of the listed states counts as
# one event. Deaths are counted automatically from the death states.
events:
  ischaemic_stroke: [AIS, AIS_AB, PB_AIS, AIS_Death, AIS_AB_Death]
  haemorrhagic_stroke: [AHS, AHS_AB, PB_AHS, AHS_Death, AHS_AB_Death]
  major_bleed: [AB, AIS_AB, AHS_AB, PS_AB, AB_Death, AIS_AB_Death, AHS_AB_Death]

# Currency symbol used in printed results and plots.
currency: "€"

# Willingness-to-pay thresholds per QALY, drawn as lines on the cost-effectiveness plane. Points
# below a line are cost-effective at that threshold. Use `wtp: []` for no lines.
wtp: [20000, 30000, 50000]

```

| Setting               | Example                   | What it does                                                                                                                      |
|-----------------------|---------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| strata                | [agegrp, sex, mm_cluster] | The stratifier columns of the input files. Every other matrix column, apart from from, is a state. [] for none                    |
| seed                  | 2026                      | The starting seed. Run $r$ uses $\text{seed} + r - 1$ , so results can be repeated exactly                                        |
| horizon               | 10                        | The number of yearly cycles, counted from cycle 0 for everyone                                                                    |
| age_limit_death_state | CurrentDeath              | The death state of patients who would pass the oldest age group                                                                   |
| discount_rate         | 0.03                      | The yearly discount rate for costs and QALYs, as a proportion (0.03 for 3%). A year in cycle $t$ counts $1 / (1 + \text{rate})^t$ |
| n_runs                | 100                       | The number of runs; each run is one point on the CE plane                                                                         |
| events                | lists of states           | The states counted as each event in the results                                                                                   |
| currency              | "€"                       | The currency symbol in the printed results and the plot                                                                           |
| wtp                   | [20000, 30000, 50000]     | Willingness-to-pay thresholds per QALY, drawn as lines on the CE plane. [] for none                                               |

`tm_read_settings()` reads the file and checks every setting. A wrong type or value stops with a message naming the setting, and a name the package does not know, such as a typo, gives a warning.

In R, the settings are a list:

```

settings <- tm_read_settings()
str(settings)

```

```

List of 9
 $ strata      : chr [1:3] "agegrp" "sex" "mm_cluster"
 $ seed       : int 2026
 $ horizon    : int 10
 $ age_limit_death_state: chr "CurrentDeath"
 $ discount_rate : num 0.03
 $ n_runs     : int 100
 $ events     :List of 3

```

```

..$ ischaemic_stroke : chr [1:5] "AIS" "AIS_AB" "PB_AIS" "AIS_Death" ...
..$ haemorrhagic_stroke: chr [1:5] "AHS" "AHS_AB" "PB_AHS" "AHS_Death" ...
..$ major_bleed : chr [1:7] "AB" "AIS_AB" "AHS_AB" "PS_AB" ...
$ currency : chr "€"
$ wtp : int [1:3] 20000 30000 50000

```

## 2 Changing settings

There are two ways:

- Edit `inputs/settings.yaml` and run `tm_run()` again. This is the way to change the analysis itself.
- Change the list in R and pass it to `tm_run()`, which is handy for trying things out; the file stays as it is.

The examples below change the list in R, with 20 runs each to keep them quick. Every call to `tm_run()` overwrites `outputs/`.

```

base <- tm_read_settings()
base$n_runs <- 20
base_results <- tm_run(base)

```

### 2.1 Discount rate

```

no_discounting <- base
no_discounting$discount_rate <- 0
no_discounting_results <- tm_run(no_discounting)

```

### 2.2 Time horizon

The example cohort enters between cycles 0 and 9. With a 5-year horizon, the patients who would enter in cycles 5 to 9 fall outside it, and `tm_run()` leaves them out with a warning:

```

five_years <- base
five_years$horizon <- 5
five_years_results <- tm_run(five_years)

```

Warning: `inputs/cohort.csv`: 1343 patients enter in cycle 5 or later, after the horizon of 5 cycles, and are left out (ids 3658, 3659, 3660, 3661, 3662, 3663, 3664, 3665, 3666, 3667 and 1333 more).

`intervention_cost.csv` lists years 1 to 9, which covers a 5-year horizon too: later years are allowed and not used.

### 2.3 Side by side

```

rbind(cbind(setting = "3% discount, 10 years", base_results$comparison),
      cbind(setting = "no discounting, 10 years", no_discounting_results$comparison),
      cbind(setting = "3% discount, 5 years", five_years_results$comparison)
)[, .(setting, delta_cost, delta_qaly, icer, quadrant)]


```

|    | setting                  | delta_cost | delta_qaly  | icer      | quadrant |
|----|--------------------------|------------|-------------|-----------|----------|
|    | <char>                   | <num>      | <num>       | <num>     | <char>   |
| 1: | 3% discount, 10 years    | -104.67869 | 0.004967325 | -21073.45 | dominant |
| 2: | no discounting, 10 years | -126.25450 | 0.005941800 | -21248.53 | dominant |
| 3: | 3% discount, 5 years     | -38.71584  | 0.001815603 | -21323.96 | dominant |

## 2.4 Willingness-to-pay lines

wtp only changes the lines on the CE plane:

```
other_lines <- base
other_lines$wtp <- c(10000, 25000)
plot(tm_run(other_lines))
```



## 2.5 Seed and number of runs

The same seed always gives the same results; another seed gives another set of runs:

```
two_runs <- base
two_runs$n_runs <- 2
first <- tm_run(two_runs)$comparison
again <- tm_run(two_runs)$comparison
two_runs$seed <- 1
other_seed <- tm_run(two_runs)$comparison
rbind(cbind(run = "seed 2026", first), cbind(run = "seed 2026 again", again),
      cbind(run = "seed 1", other_seed))[, .(run, delta_cost, delta_qaly)]
```

```

        run delta_cost  delta_qaly
      <char>      <num>      <num>
1:      seed 2026  -102.4714  0.003580796
2: seed 2026 again  -102.4714  0.003580796
3:      seed 1    -118.5145  0.004084200

```

More runs give more points on the CE plane and a steadier mean. Each run of the 5,000 example patients takes about a quarter of a second.